


## Brain Tumor Classification (CT Scan Images) using SkipNet Architecture



```
import numpy as np
import pandas as pd

base_path = "/kaggle/input/brain-tumor-multimodal-image-ct-and-
mri/Dataset/Brain Tumor CT scan Images/"
categories = ["Healthy", "Tumor"]

image_paths = []
labels = []

for category in categories:
    category_path = os.path.join(base_path, category)
    for image_name in os.listdir(category_path):
        image_path = os.path.join(category_path, image_name)
        image_paths.append(image_path)
        labels.append(category)

df = pd.DataFrame({
    "image_path": image_paths,
    "label": labels
})
```

```
df.head()
```

	image_path	label
0	/kaggle/input/brain-tumor-multimodal-image-ct-...	Healthy
1	/kaggle/input/brain-tumor-multimodal-image-ct-...	Healthy
2	/kaggle/input/brain-tumor-multimodal-image-ct-...	Healthy
3	/kaggle/input/brain-tumor-multimodal-image-ct-...	Healthy
4	/kaggle/input/brain-tumor-multimodal-image-ct-...	Healthy

```
df.tail()
```

	image_path	label
4613	/kaggle/input/brain-tumor-multimodal-image-ct-...	Tumor
4614	/kaggle/input/brain-tumor-multimodal-image-ct-...	Tumor
4615	/kaggle/input/brain-tumor-multimodal-image-ct-...	Tumor
4616	/kaggle/input/brain-tumor-multimodal-image-ct-...	Tumor
4617	/kaggle/input/brain-tumor-multimodal-image-ct-...	Tumor

```
df.shape
```

```
(4618, 2)
```

```
df.columns
```

```
Index(['image_path', 'label'], dtype='object')
```

```
df.duplicated().sum()
```

```
0
```

```
df.isnull().sum()
```

```
image_path    0
label         0
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4618 entries, 0 to 4617
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   image_path  4618 non-null   object
 1   label       4618 non-null   object
dtypes: object(2)
memory usage: 72.3+ KB
```

```
df['label'].unique()
```

```
array(['Healthy', 'Tumor'], dtype=object)
```

```
df['label'].value_counts()
```

```

label
Tumor      2318
Healthy     2300
Name: count, dtype: int64

import seaborn as sns
import matplotlib.pyplot as plt


plt.figure(figsize=(8, 6))
sns.countplot(data=df, x="label", palette="viridis")
plt.title("Distribution of Labels - Count Plot")
plt.xlabel("Tumor Type")
plt.ylabel("Count")

for p in plt.gca().patches:
    plt.gca().annotate(f'{{int(p.get_height())}}',
                      (p.get_x() + p.get_width() / 2., p.get_height()),
                      ha='center', va='center', fontsize=11, color='black',
xytext=(0, 5),
                      textcoords='offset points')

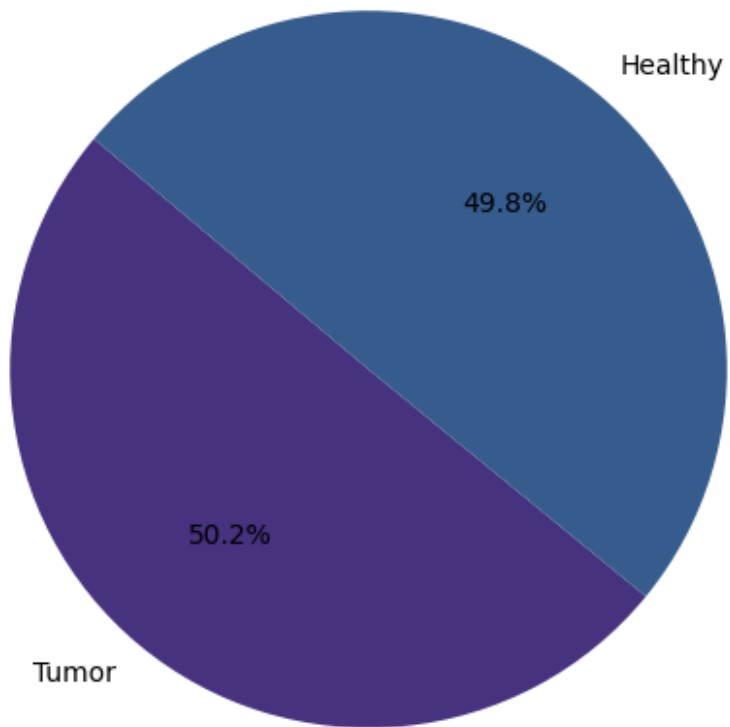
plt.show()

label_counts = df['label'].value_counts()
plt.figure(figsize=(8, 6))
plt.pie(label_counts, labels=label_counts.index, autopct='%1.1f%%',
startangle=140, colors=sns.color_palette("viridis"))
plt.title("Distribution of Labels - Pie Chart")
plt.show()

```



Distribution of Labels - Pie Chart



```
import cv2

num_images = 5

plt.figure(figsize=(15, 12))


for i, category in enumerate(categories):
    category_images = df[df['label'] ==
category]['image_path'].iloc[:num_images]

    for j, img_path in enumerate(category_images):

        img = cv2.imread(img_path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

        plt.subplot(len(categories), num_images, i * num_images + j + 1)
        plt.imshow(img)
        plt.axis('off')
        plt.title(category)
```

```
plt.tight_layout()
plt.show()
```



```
from sklearn.preprocessing import LabelEncoder
```

```
label_encoder = LabelEncoder()
```

```
df['category_encoded'] = label_encoder.fit_transform(df['label'])
```

```
df = df[['image_path', 'category_encoded']]
```

```
from imblearn.over_sampling import RandomOverSampler
```

```
ros = RandomOverSampler(random_state=42)
```

```
X_resampled, y_resampled = ros.fit_resample(df[['image_path']],
df['category_encoded'])
```

```
df_resampled = pd.DataFrame(X_resampled, columns=['image_path'])
```

```
df_resampled['category_encoded'] = y_resampled
```

```
print("\nClass distribution after oversampling:")
```

```
print(df_resampled['category_encoded'].value_counts())
```

```
Class distribution after oversampling:
```

```
category_encoded
```

```
0    2318
```

```
1    2318
```

```
Name: count, dtype: int64
```

```
df_resampled
```

	image_path	category_encoded
0	/kaggle/input/brain-tumor-multimodal-image-ct-...	0
1	/kaggle/input/brain-tumor-multimodal-image-ct-...	0
2	/kaggle/input/brain-tumor-multimodal-image-ct-...	0
3	/kaggle/input/brain-tumor-multimodal-image-ct-...	0
4	/kaggle/input/brain-tumor-multimodal-image-ct-...	0
...	...	...
4631	/kaggle/input/brain-tumor-multimodal-image-ct-...	0
4632	/kaggle/input/brain-tumor-multimodal-image-ct-...	0
4633	/kaggle/input/brain-tumor-multimodal-image-ct-...	0
4634	/kaggle/input/brain-tumor-multimodal-image-ct-...	0
4635	/kaggle/input/brain-tumor-multimodal-image-ct-...	0

[4636 rows x 2 columns]

```
df_resampled['category_encoded'] =
df_resampled['category_encoded'].astype(str)
```

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
```

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
Activation, Dropout, BatchNormalization
from tensorflow.keras import regularizers
```

```
import warnings
warnings.filterwarnings("ignore")
```

```
print ('check')
```

check

```
train_df_new, temp_df_new = train_test_split(
    df_resampled,
    train_size=0.8,
    shuffle=True,
    random_state=42,
    stratify=df_resampled['category_encoded']
)
```

```
valid_df_new, test_df_new = train_test_split(
    temp_df_new,
    test_size=0.5,
    shuffle=True,
    random_state=42,
```

```

        stratify=temp_df_new['category_encoded']
    )

    batch_size = 16
    img_size = (224, 224)
    channels = 3
    img_shape = (img_size[0], img_size[1], channels)

    tr_gen = ImageDataGenerator(rescale=1./255)
    ts_gen = ImageDataGenerator(rescale=1./255)

    train_gen_new = tr_gen.flow_from_dataframe(
        train_df_new,
        x_col='image_path',
        y_col='category_encoded',
        target_size=img_size,
        class_mode='binary',
        color_mode='rgb',
        shuffle=True,
        batch_size=batch_size
    )

    valid_gen_new = ts_gen.flow_from_dataframe(
        valid_df_new,
        x_col='image_path',
        y_col='category_encoded',
        target_size=img_size,
        class_mode='binary',
        color_mode='rgb',
        shuffle=True,
        batch_size=batch_size
    )

    test_gen_new = ts_gen.flow_from_dataframe(
        test_df_new,
        x_col='image_path',
        y_col='category_encoded',
        target_size=img_size,
        class_mode='binary',
        color_mode='rgb',
        shuffle=False,
        batch_size=batch_size
    )

```

Found 3708 validated image filenames belonging to 2 classes.  
 Found 464 validated image filenames belonging to 2 classes.  
 Found 464 validated image filenames belonging to 2 classes.



```

import tensorflow as tf

print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))

Num GPUs Available:  2

gpus = tf.config.list_physical_devices('GPU')
if gpus:
    try:
        for gpu in gpus:
            tf.config.experimental.set_memory_growth(gpu, True)
        print("GPU is set for TensorFlow")
    except RuntimeError as e:
        print(e)

GPU is set for TensorFlow

import torch
import torch.nn as nn
import torch.nn.functional as F

class SpatialAttentionLayer(nn.Module):
    def __init__(self, in_channels):
        super(SpatialAttentionLayer, self).__init__()
        self.conv1 = nn.Conv2d(in_channels, in_channels // 2, kernel_size=1)
        self.conv2 = nn.Conv2d(in_channels // 2, 1, kernel_size=3, padding=1,
dilation=1)
        self.conv3 = nn.Conv2d(1, 1, kernel_size=3, padding=2, dilation=2)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        attention = self.conv1(x)
        attention = F.relu(self.conv2(attention))
        attention = self.sigmoid(self.conv3(attention))
        return x * attention

class CNBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(CNBlock, self).__init__()
        self.spatial_attention = SpatialAttentionLayer(in_channels)
        self.conv = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, kernel_size=3, padding=1),
            nn.BatchNorm2d(out_channels),
            nn.ReLU()
        )

    def forward(self, x):
        x = self.spatial_attention(x)
        return self.conv(x)

```

```

class SKIPNet(nn.Module):
    def __init__(self, num_classes=3):
        super(SKIPNet, self).__init__()
        self.block1 = CNNBlock(3, 32)
        self.block2 = CNNBlock(32, 64)
        self.block3 = CNNBlock(64, 128)
        self.dropout = nn.Dropout(0.5)
        self.downsample = nn.Conv2d(128, 128, kernel_size=2, stride=2)
        self.final_attention = SpatialAttentionLayer(128)
        self.fc = nn.Sequential(
            nn.Flatten(),
            nn.Linear(128 * 16 * 16, 256),
            nn.ReLU(),
            nn.Linear(256, num_classes)
        )

    def forward(self, x):
        x = self.block1(x)
        x = self.dropout(x)
        x = self.block2(x)
        x = self.dropout(x)
        x = self.block3(x)
        x = self.dropout(x)
        x = self.downsample(x)
        x = self.final_attention(x)
        x = self.fc(x)
        return x

model = SKIPNet(num_classes=3)
print(model)

SKIPNet(
  (block1): CNNBlock(
    (spatial_attention): SpatialAttentionLayer(
      (conv1): Conv2d(3, 1, kernel_size=(1, 1), stride=(1, 1))
      (conv2): Conv2d(1, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      (conv3): Conv2d(1, 1, kernel_size=(3, 3), stride=(1, 1), padding=(2,
2), dilation=(2, 2))
      (sigmoid): Sigmoid()
    )
    (conv): Sequential(
      (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): ReLU()
    )
  )
  (block2): CNNBlock(
    (spatial_attention): SpatialAttentionLayer(

```

```

        (conv1): Conv2d(32, 16, kernel_size=(1, 1), stride=(1, 1))
        (conv2): Conv2d(16, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
        (conv3): Conv2d(1, 1, kernel_size=(3, 3), stride=(1, 1), padding=(2,
2), dilation=(2, 2))
        (sigmoid): Sigmoid()
    )
    (conv): Sequential(
      (0): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): ReLU()
    )
  )
  (block3): CNNBlock(
    (spatial_attention): SpatialAttentionLayer(
      (conv1): Conv2d(64, 32, kernel_size=(1, 1), stride=(1, 1))
      (conv2): Conv2d(32, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
      (conv3): Conv2d(1, 1, kernel_size=(3, 3), stride=(1, 1), padding=(2,
2), dilation=(2, 2))
      (sigmoid): Sigmoid()
    )
    (conv): Sequential(
      (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): ReLU()
    )
  )
  (dropout): Dropout(p=0.5, inplace=False)
  (downsample): Conv2d(128, 128, kernel_size=(2, 2), stride=(2, 2))
  (final_attention): SpatialAttentionLayer(
    (conv1): Conv2d(128, 64, kernel_size=(1, 1), stride=(1, 1))
    (conv2): Conv2d(64, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv3): Conv2d(1, 1, kernel_size=(3, 3), stride=(1, 1), padding=(2, 2),
dilation=(2, 2))
    (sigmoid): Sigmoid()
  )
  (fc): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
    (1): Linear(in_features=32768, out_features=256, bias=True)
    (2): ReLU()
    (3): Linear(in_features=256, out_features=3, bias=True)
  )
)

```

```

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

```

```

class SpatialAttentionLayer(keras.layers.Layer):
    def __init__(self, **kwargs):
        super(SpatialAttentionLayer, self).__init__(**kwargs)
        self.conv1 = layers.Conv2D(32, kernel_size=1, activation='relu')
        self.conv2 = layers.Conv2D(1, kernel_size=3, padding='same',
activation='sigmoid')

    def call(self, inputs):
        attention = self.conv1(inputs)
        attention = self.conv2(attention)
        return inputs * attention

class CNNBBlock(keras.layers.Layer):
    def __init__(self, filters, **kwargs):
        super(CNNBlock, self).__init__(**kwargs)
        self.spatial_attention = SpatialAttentionLayer()
        self.conv = keras.Sequential([
            layers.Conv2D(filters, kernel_size=3, padding='same',
activation='relu'),
            layers.BatchNormalization(),
            layers.ReLU()
        ])

    def call(self, inputs):
        x = self.spatial_attention(inputs)
        return self.conv(x)

class SKIPNet(keras.Model):
    def __init__(self, num_classes=2):
        super(SKIPNet, self).__init__()
        self.block1 = CNNBBlock(32)
        self.block2 = CNNBBlock(64)
        self.block3 = CNNBBlock(128)
        self.dropout = layers.Dropout(0.5)
        self.downsample = layers.Conv2D(128, kernel_size=2, strides=2)
        self.final_attention = SpatialAttentionLayer()
        self.flatten = layers.Flatten()
        self.fc1 = layers.Dense(256, activation='relu')
        self.fc2 = layers.Dense(1, activation='sigmoid')

    def call(self, inputs):
        x = self.block1(inputs)
        x = self.dropout(x)
        x = self.block2(x)
        x = self.dropout(x)
        x = self.block3(x)
        x = self.dropout(x)
        x = self.downsample(x)

```

```

        x = self.final_attention(x)
        x = self.flatten(x)
        x = self.fc1(x)
        return self.fc2(x)

model = SKIPNet(num_classes=2)
model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])

model.build(input_shape=(None, 224, 224, 3))

history = model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=5,
    batch_size=16,
    verbose=1
)

Epoch 1/5
232/232 _____ 143s 400ms/step - accuracy: 0.7252 - loss:
0.7376 - val_accuracy: 0.9224 - val_loss: 0.2028
Epoch 2/5
232/232 _____ 68s 291ms/step - accuracy: 0.9617 - loss: 0.1009
- val_accuracy: 0.9116 - val_loss: 0.2602
Epoch 3/5
232/232 _____ 68s 292ms/step - accuracy: 0.9659 - loss: 0.1051
- val_accuracy: 0.9483 - val_loss: 0.1678
Epoch 4/5
232/232 _____ 68s 292ms/step - accuracy: 0.9895 - loss: 0.0318
- val_accuracy: 0.9591 - val_loss: 0.2458
Epoch 5/5
232/232 _____ 68s 292ms/step - accuracy: 0.9906 - loss: 0.0217
- val_accuracy: 0.9440 - val_loss: 0.3404

from sklearn.metrics import confusion_matrix, classification_report

def plot_history(history):
    plt.figure(figsize=(12, 4))


    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.title('Training and Validation Accuracy')

    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'], label='Train Loss')

```

```
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Training and Validation Loss')
plt.show()
```


```
plot_history(history)
```



```
y_pred = model.predict(test_gen_new)
y_pred_classes = (y_pred > 0.5).astype(int).flatten()
y_true = test_gen_new.classes
```

29/29 ————— 5s 149ms/step

```
conf_matrix = confusion_matrix(y_true, y_pred_classes)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()
```



```
print('Classification Report:\n', classification_report(y_true,  
y_pred_classes))
```

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.97	0.95	232
1	0.96	0.94	0.95	232
accuracy			0.95	464
macro avg	0.95	0.95	0.95	464
weighted avg	0.95	0.95	0.95	464